

Meeting Briefing - Magdalena (Sales) & Zuza (Customer Engineer) - Google Cloud

Time: April 17, 2026 @ 10:00 AM (online)

Duration: Assume 30-45 min

Goal: Qualify for AI Tier of Google for Startups Cloud Program + explore workload expansion

A. Your 30-Second Pitch (open with this)

"We're Spectre AI - the intelligence layer for financial research and trading. We unify AI-powered research, social intelligence, and execution in one workspace across stocks, crypto, and commodities. Users talk to data instead of hunting through dashboards. We've been on Google Cloud for 2 years through the Web3 program, and we've now evolved into a full AI-native company - which is why we're re-applying under the AI tier."

B. Current GCP Footprint (March 2026 bill: ~z■3,446 / ~\$870)

Service	Spend (PLN)	What it does	Growth signal
Firebase Realtime Database	1,619	Live data sync (prices, agent state, activity log)	+3720% MoM - scaling fast
Cloud Run	1,076	2 services: appresearchbeta (market intel backend) + x-dash-api (X/Twitter intelligence, EU-West1)	+312%
Artifact Registry	310	Docker images for Cloud Run	+202%
Cloud Run Functions	304	SearchEngineApiV4 (search abstraction)	+304%
Compute Engine	247	Supporting VMs	+83%
Gemini API	23	Minimal - early testing only	+23%
Secret Manager, Cloud Storage, Networking	~115	Supporting	-

Key story: usage has **exploded 3-37x MoM** across core services. You are ramping. This is the single most important slide in the meeting.

C. Honest Reality Check (things NOT to overclaim)

The earlier email said "100+ serverless functions on GCP" and "510 endpoints." The actual picture:

- **2 Cloud Run services on GCP** (appresearchbeta, x-dash-api) + 1 Cloud Function (SearchEngineApiV4)
- **~120 serverless endpoints on Vercel** (not GCP) - handle the research/trading apps
- **19 AI agents in packages/server/agents/** - all powered by Perplexity Sonar Pro, Claude, OpenAI (NOT Gemini currently)

If Zuza probes the architecture, be direct: "Our serverless frontend API layer is on Vercel today. Our heavy compute (market intelligence backend + X/Twitter scraping pipeline) runs on Cloud Run. That's a split we may want to revisit with your team."

D. Architecture Snapshot (for Zuza's technical questions)

Browser (apps/research :5180, apps/trading :5181)

```
|
+--> Vercel serverless (120 endpoints, /api/*) <-- frontend API layer
|
+--> /ext-api/* --> Cloud Run: appresearchbeta (US-Central1)
|                     (charts, liquidations, token search, sector analysis)
|
+--> /spectre-api/* --> Cloud Functions: SearchEngineApiV4
|
+--> X-Dash data --> Cloud Run: x-dash-api (EU-West1)
|                     (KOL mentions, narrative tracking, sentiment)
```

Server-side (Express, packages/server/) -- local dev; serverless in prod

```
+--> 19 AI Agents --> Perplexity Sonar Pro (primary), Claude, OpenAI
+--> Codex GraphQL (token 0HLCV, 13K+ tokens)
+--> Binance klines (major tokens)
+--> Firebase Realtime Database (live price sync, agent state)
```

Tech stack: React 18 + Vite, Express, Node.js, plain CSS. No Kubernetes. No BigQuery. No Pub/Sub. Caching on Vercel KV + in-memory (TTLs 30s-10min).

E. The Ask - What You Want from This Meeting

1. AI Tier enrollment in Google for Startups Cloud Program

- You've submitted the form. Confirm status, timeline, credit amount.

2. Committed-Use Discounts (CUDs) on Cloud Run + Firebase RTDB

- Spend is scaling 3-37x MoM. Discuss 1y/3y commitment pricing.

3. Malaga Security Event recordings (can't attend but interested in materials)

4. Intro to Google for Startups program team (if not already via Magdalena)

F. Google's Likely Pitches (and your positioning)

F1. "Migrate to Vertex AI from AI Studio"

- **Your line:** "Gemini is ~\$23/month today. Most of our AI is Perplexity Sonar Pro (news generation, research articles) and Claude (Monarch agent). We're happy to evaluate Vertex AI for future Gemini workloads - specifically if it gives us lower latency and dedicated capacity for financial content generation."
- **Probe back:** Does Vertex AI support bringing our own Perplexity/Claude via Model Garden or partnerships? What's the real cost vs latency tradeoff vs direct API?

F2. "Gemini Enterprise for agent orchestration"

- **Your line:** "We have 19 production agents today - news curator, daily brief, breaking news, entity resolver, token analysis, calendar analysis. They run on schedulers with a rate-limited queue (`perplexityGate.js`). Show us what Gemini Enterprise adds over our current LangChain-style orchestration."
- **Probe:** What does Gemini Enterprise cost? Is there a free tier for AI Program members?

F3. "BigQuery for analytics"

- **Your line:** "We don't have an analytics warehouse today. We stream prices, track 13K+ tokens, log agent activity. If we need historical analytics (backtesting strategies, user behavior trends), BigQuery is a strong fit. Not urgent - Q3 2026 timeline."
- **Probe:** Any free credits for BigQuery through the AI Program? Data transfer costs from Cloud Run?

F4. "Move Vercel workloads to Cloud Run"

- **Your line:** "We're open to it. Vercel is good for our React/Vite frontend apps but we'd consider consolidating the ~120 serverless endpoints onto Cloud Run if the economics + DX work. What's the migration tooling?"
- **Probe:** Performance benchmarks for cold starts? Cloud Run Functions vs Vercel Functions? Any case studies?

G. Technical Questions Zuza May Ask

Q	Your answer
"What's driving the 3720% Firebase growth?"	Live price streams for 13K+ tokens, agent state sync, real-time activity feed (200-event log)
"Why Perplexity over Gemini?"	Sonar Pro has lowest latency for our use case (live news generation with web search built-in). Happy to benchmark Gemini 2.5 for new workloads.
"How many concurrent users?"	[YOU FILL - check analytics]
"What's your data volume / QPS?"	Codex queries ~50-100 qps peak, Binance klines ~20 qps, news RSS every 5 min
"Multi-region strategy?"	Currently US-Central1 + Europe-West1 (x-dash-api). Most users in US/EU.
"DR / backups?"	Relying on Vercel + Firebase managed replication. No formal DR plan yet.
"Compliance / SOC2?"	Not yet. On roadmap for enterprise customer readiness.
"Why not Vertex AI already?"	Minimal Gemini usage so far - AI Studio was fine for sandbox. As we scale Gemini, Vertex AI migration is logical.

H. Workloads You Can Genuinely Expand on GCP

Pitch these as "what we're thinking about next" - it shows growth intent without overcommitting:

1. **Vertex AI for a new feature** - e.g., a Gemini-powered "chat with your portfolio" agent (competes with Monarch). Good pilot workload.
2. **Firestore / Cloud SQL** for user profiles + watchlists (currently Vercel KV). Consolidation story.
3. **BigQuery** for historical price data + user behavior analytics (Q3-Q4 2026).
4. **Cloud CDN / Cloud Storage** for token logos, brief audio, generated hero images (currently mixed hosting).
5. **Pub/Sub** for the news pipeline fan-out (curator -> writer -> publisher).

I. Red Flags to Avoid

- **Don't oversell Gemini usage.** Magdalena already knows it's ~23. Be honest.
- **Don't claim "100+ functions on GCP"** when it's 2 Cloud Run services. Zuza will check.
- **Don't commit to migration timelines** without Sunny + team alignment.
- **Don't mention competitors by name** (Vercel, Perplexity, Anthropic) in a negative way - just as current tools.

J. Meeting Close

End with:

1. **Clear next step:** "Can you confirm AI Tier application status by next week?"
 2. **Ask for the CE follow-up:** "Zuza, can you send us a short architecture recommendation doc based on today?"
 3. **Offer mutual value:** "If AI Program approves, we'd love to be a case study for AI-native fintech on GCP."
-

K. Files Referenced for This Briefing

- `packages/server/index.js` - Cloud Run endpoints
- `packages/server/agents/` - 19 AI agents
- `apps/research/vercel.json` - GCP rewrites (lines 132-133)
- `.env.example` - X_DASH_API config
- `docs/spectre-research-api.md` - appresearchbeta architecture